Package Vignette for ndtv: Network Dynamic Temporal Visualizations (Version 0.6)

Skye Bender-deMoll

April 21, 2015

Contents

1 Introduction

The Network Dynamic Temporal Visualization (`ndtv`) package provides tools
for visualizing changes in network structure and attributes over time. It works
with dynamic network information encoded in `networkDynamic` (Butts et al.
, 2015) objects as its input, and outputs animated movies. The package will
eventually include timelines and other types of dynamic visualizations of evolv-
ing relational structures. The core use-case for development is examining the
output of statistical network models (such as those produced by the `tergm`
(Krivitsky et al. , 2014) package in `statnet` (Handcock et al , 2003)) and sim-
ulations of disease spread across networks. The `ndtv` (Bender-deMoll , 2015)
package relies on many other packages to do much of the heavy lifting, especially
`animation` (Yihui, Xie et al. , 2013) and `networkDynamic` and requires exter-
nal libraries (FFmpeg) to save movies out of the R environment. To use `ndtv`
effectively you must be already familiar with the functionality and assumptions
of `networkDynamic`. This package is intended to eventually replace much of the
functionality in the `rSoNIA` package (Bender-deMoll et al., 2008).

This work was supported by grant R01HD68395 from the National Institute
of Health.

2 A quick example

2.1 Reinventing the wheel

Lets get started! We can render a trivially simple animation in the R plot
window.

```
> library(ndtv) # also loads animation and networkDynamic
> wheel <- network.initialize(10)  # create a toy network
> add.edges.active(wheel,tail=1:9,head=c(2:9,1),onset=1:9, terminus=11)
> add.edges.active(wheel,tail=10,head=c(1:9),onset=10, terminus=12)
> plot(wheel) # peek at the static version
> render.animation(wheel) # compute and render

slice parameters:
  start:1
  end:12
  interval:1
  aggregate.dur:1
  rule:latest

> ani.replay() # play back in plot window
```

Hopefully, when you ran `ani.replay()` you saw a bunch of labeled nodes moving smoothly around in the R plot window, with edges slowly appearing to link them into a circle. Finally a set of "spoke" edges appear to draw a vertex into the center. If that didn't work, the footnote has a link to an example of the movie [1] you are supposed to see. For some kinds of networks the animated version gives a very different impression of the connectivity of the network that a static plot of the same network (Figure 1)

2.2   What just happened?

Simple right? Yes, but that is because most of the difficult parts happened under the hood using default values. In a nutshell, this is how it worked:

1. We created a `networkDynamic` object named `wheel` containing information about the timing of edge activity.

2. `render.animation()` asked the package to create an animation for `wheel` but we didn't include any arguments indicating what should be rendered or how.

3. Since `render.animation()` didn't find any stored coordinate information about where to draw the vertices and edges, it (invisibly) called `compute.animation()` with default arguments to figure out where to position the vertices at each time step.

4. Because we didn't tell `compute.animation()` what time points to look at when doing its computations, it reported this, `"No slice.par found"`, and made a guess as to when the animation should start and end (the earliest and latest observed times in the network) and how much time should be incremented between each set of layout coordinate calculations.

[1] http://statnet.csde.washington.edu/movies/ndtv_vignette/wheel.mp4

Figure 1: Standard network plot of our trivial "wheel" network does not reveal dynamics. Compare with animated movie version: `http://statnet.csde.washington.edu/movies/ndtv_vignette/wheel.mp4`

5. `compute.animation()` then stepped through the `wheel` network, computing coordinates for each time step and storing them. (This was the `"Calculating layout for network slice from time 1 to 2"` ... part.)

6. `render.animation()` also stepped through the network, using the stored coordinates, `plot.network()` and `ani.record()` functions to cache snapshots of the network. It also caches a number of "tweening" images between each time step to smoothly interpolate the positions of the vertices. `"rendering 10 frames for slice 1"` ...

7. `ani.replay()` quickly redraws the sequence of cached images in the plot window as an animation.

Of course, using defaults doesn't give much control of what should be rendered and how it should look. For more precise control of the processes, layout algorithms, etc, we can call each of the steps in sequence.

3   A tergm simulation example

Lets look at a more realistic example using output from the simulation of a crude dynamic model. This uses the statnet `tergm` package to estimate the parameters for an edge formation and dissolution process which produces a network similar to the Florentine business network (`?ergm::flobusiness`) given as input. Once the model has been estimated, we can take a number of sequential draws from it to see how the network might "evolve" over time. When we generate the movie, we can include the model statistics on screen to see how they are influenced by edge additions and deletions. This example also assumes you have some of the external libraries working (Java and FFmpeg) so you run into problems, try skipping to Dependencies (section 9) and come back.

If you are not interested in running the model, you can just load a corresponding example data object with `data(stergm.sim.1)`.
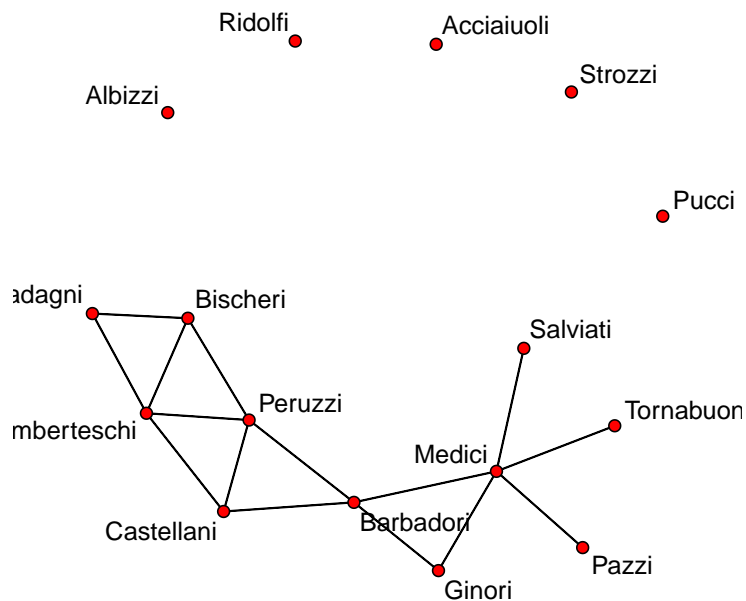
3.1   Data Setup

First load in the main necessary libraries (each of which loads a bunch of additional R libraries).

```
> require(ndtv)     # dynamic network animations
> require(tergm)    # dynamic ergm simulations
```

Load in the original Florentine business network.

```
> data("florentine") # an example network
> plot(flobusiness,displaylabels=TRUE)
```

Define basic `stergm` model with formation and dissolution parameters.

```
> theta.diss <- log(9)
> stergm.fit.1 <- stergm(flobusiness,
+    formation= ~edges+gwesp(0,fixed=TRUE),
+    dissolution = ~offset(edges),
+        targets="formation",
+    offset.coef.diss = theta.diss,
+    estimate = "EGMME"         )
```

(time passes, lots simulation status output hidden)

Now we can simulate 100 discrete time steps from the model and save them as a `dynamicNetwork` object.

```
> stergm.sim.1 <- simulate.stergm(stergm.fit.1,
+                     nsim=1, time.slices = 100)
```

3.2 Animation Setup

Since this isn't a terribly exciting simulation, lets only calculate coordinates for part of the simulated time period by using the `start` and `end` parameters of `slice.par` to specify a time range. We can also ask it to use the MDSJ layout (assuming MDSJ and Java are installed).

```
> slice.par<-list(start=75,end=100,interval=1,
+                 aggregate.dur=1,rule="latest")
> compute.animation(stergm.sim.1,slice.par=slice.par,
+                                  animation.mode='MDSJ')
```

3.3 Playing an animation in R plot window

Now that we have all the coordinates stored, we can define some parameters for `render.par` to specify how many `tween.frames` to render, and tell it to display the time and the summary statistics formula.

```
> render.par=list(tween.frames=5,show.time=TRUE,
+                 show.stats="~edges+gwesp(0,fixed=TRUE)")
```

Then we ask it to graphically render the animation, passing in some of the standard `plot.network` graphics arguments to change the color of the edges and show the labels with a smaller size and blue color.

```
> render.animation(stergm.sim.1,render.par=render.par,
+                  edge.col="darkgray",displaylabels=TRUE,
+                  label.cex=.6,label.col="blue")
```

This takes some time and produces many lines output which we are not showing. The output could also be suppressed by adding a `verbose=FALSE` argument.

After it has finished, replay the movie in an R plot window.

```
> ani.replay()
```

Notice that in addition to the labels on the bottom of the plot indicating which time step is being viewed, it also displays the network statistics of interest for the time step. When the "edges" parameter increases up, you can see the density on the graph increase and the number of isolates decrease. Eventually the model corrects, and the parameter drifts back down.

3.4 Saving an animation as video

We can also use the `animation` library to save out the movie in `.mp4` format (assuming that the FFmpeg library is installed on your machine).

```
> saveVideo(ani.replay(),video.name="stergm.sim.1.mp4",
+                  other.opts="-b 5000k",clean=TRUE)
```

```
NULL
```

This should produce a movie[2] in an R working directory on disk. The `other.opts` parameter is set here to generate a higher-quality video than the default, but this will result in a large file size. For more information on compressing videos for the web, see Compressing Video (section 10).

3.5   Viewing animation as a web page

An alternate way to view the animation is to render it out as an HTML5 animation embedded in a web page using the ndtv-d3 player (Michalec, G., et al. , 2014). This has the advantage of not requiring any external libraries, but does need a modern web browser to display properly. `render.d3movie` operates similarly to `render.animation`, using the same coordinates stored by `compute.animation`.

```
> render.d3movie(stergm.sim.1,render.par=render.par,
+                edge.col="darkgray",displaylabels=TRUE,
+                label.cex=.6,label.col="blue",
+                filename='stergm.sim1.html')
```
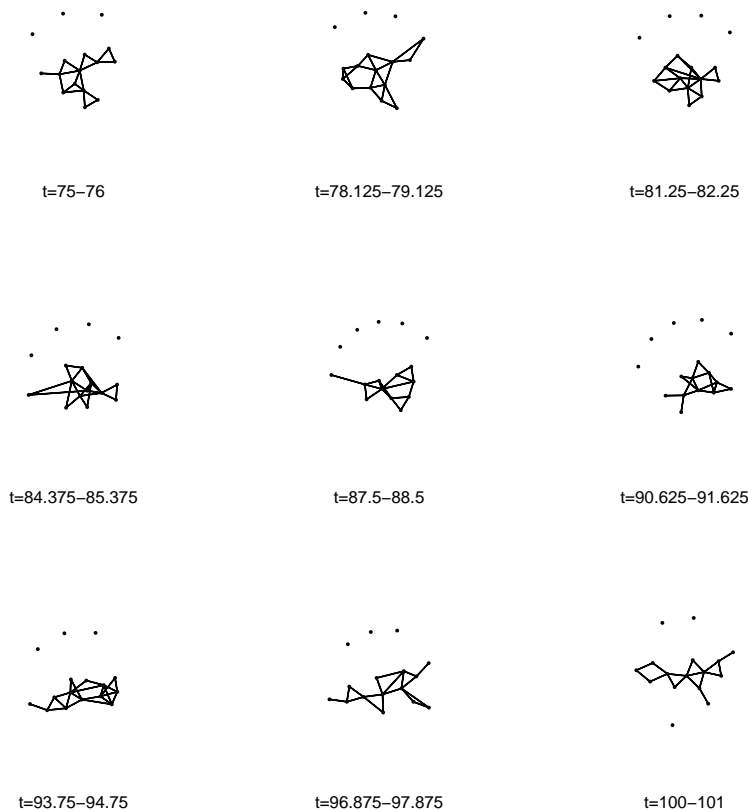
This should open web browser displaying the page with a Javascript animation player and an interactive version of the network. There is an additional vignette demoing the features (such as interactive tooltips, and embedding in Rmarkdown documents) of the ndtv-d3 player included in the ndtv package. For an example, see `http://statnet.github.io/ndtv-d3/`

3.6   Other views

There is also a `filmstrip()` function that will create a "small multiple" plot using frames of the animation to construct a visual summary of the network changes as a static plot.

```
> filmstrip(stergm.sim.1,displaylabels=FALSE)
```

---

[2]`http://statnet.csde.washington.edu/movies/ndtv_vignette/stergm.sim.1.mp4`

t=75–76 t=78.125–79.125 t=81.25–82.25

t=84.375–85.375 t=87.5–88.5 t=90.625–91.625

t=93.75–94.75 t=96.875–97.875 t=100–101

4   Slicing time

The basic network layout algorithms we are using, like most "traditional' network metrics, don't really know what to do with dynamic networks. They need to be fed a static set of relationships which can be used to compute a set of distances in a Euclidean space suitable for plotting. A common way to apply static metrics to a time-varying object is to sample it, taking a sequence static observations at a series of time points and using these to describe the changes over time. In the case of networks, we call this "extracting" or "slicing".

Slicing up a dynamic network created from discrete panels may be fairly straightforward but it is much less clear how to do it when working with continuous time or streaming relations. How often should we slice? Should the slices measure the state of the network at a specific instant, or aggregate over a longer time period? The answer probably depends on what the important features to visualize are in your data-set. The `slice.par` parameters make it possible to

experiment with various slicing options. In many situations we have even found (Bender-deMoll and McFarland , 2006) it useful to let slices mostly overlap – incrementing each one by a small value to help show fluid changes on a moderate timescale instead of the rapid changes happening on a fine timescale.

As an example, lets look at the McFarland (McFarland , 2001) data-set of streaming classroom interactions and see what happens when we chop it up in various ways. First, we can animate at the fine time scale, viewing the first half-hour of class using instantaneous slices.

```
> data(McFarland_cls33_10_16_96)
> slice.par<-list(start=0,end=30,interval=2.5,
+                 aggregate.dur=0,rule="latest")
> compute.animation(cls33_10_16_96,
+                 slice.par=slice.par,animation.mode='MDSJ')
> render.animation(cls33_10_16_96,
+                  displaylabels=FALSE,vertex.cex=1.5)
> ani.replay()
```

We can also get an idea of how we are slicing up the network by using the `timeline()` function to plot the `slice.par` parameters against the vertex and edge spells. Our very thin slices (gray vertical lines) (`aggregate.dur=0`) are not intersecting many edge events (purple numbers) at once.

```
> timeline(cls33_10_16_96,slice.par=slice.par)
```

Notice that in the animation most of the vertices are isolates, occasionally linked into brief pairs or stars by speech acts[3]. However, if we aggregate over a longer time period of 2.5 minutes we start to see the individual acts form into triads and groups[4].

```
> slice.par<-list(start=0,end=30,interval=2.5,
+                 aggregate.dur=2.5,rule="latest")
> compute.animation(cls33_10_16_96,
+                 slice.par=slice.par,animation.mode='MDSJ')
> render.animation(cls33_10_16_96,
+                  displaylabels=FALSE,vertex.cex=1.5)
> ani.replay()
```

To reveal slower structural patterns we can make the aggregation period even longer, and let the slices overlap (by making `interval` less than `aggregate.dur`)

---

[3]http://statnet.csde.washington.edu/movies/ndtv_vignette/cls33_10_16_96v1.mp4
[4]http://statnet.csde.washington.edu/movies/ndtv_vignette/cls33_10_16_96v2.mp4

so that the changes will be less dramatic between successive views[5].

```
> slice.par<-list(start=0,end=30,interval=1,
+                 aggregate.dur=5,rule="latest")
> timeline(cls33_10_16_96,slice.par=slice.par)
> compute.animation(cls33_10_16_96,
+                 slice.par=slice.par,animation.mode='MDSJ')
> render.animation(cls33_10_16_96,
+                  displaylabels=FALSE,vertex.cex=1.5)
> ani.replay()
```


Note that when we use a long duration slice, it is quite likely that the edge between a pair of vertices has more than one active period. How should this condition be handled? If the edge has attributes, which ones should be shown? Ideally we might want to aggregate the edges in some way, perhaps adding the weights together. Currently edge attributes are not aggregated and the `rule`

---

[5]http://statnet.csde.washington.edu/movies/ndtv_vignette/cls33_10_16_96v3.mp4

element of the `slice.par` argument controls which attribute should be returned for an edge when multiple elements are encountered. Generally `rule='latest'` gives reasonable results, returning the most recent value found within the query spell.

5   Layout algorithms for animations

Producing "good" (for an admittedly ambiguous definition of good) layouts of networks is generally a computationally difficult problem. There are a wide variety of algorithms and approaches being developed. Doing layouts for animations adds additional challenges because it is usually desirable that the layouts remain stable over time. Ideally this means that the layouts don't change much unless the network structure changes, and that small changes in the network structure should lead to small changes in the layouts. Many otherwise excellent static layout algorithms are not stable in this sense, or they may require very specific parameter settings to improve their results for animation applications.

The `network.layout.animate.*` layouts included in `ndtv` are adaptations or wrappers for existing static layout algorithms with some appropriate parameter presets. They all accept the coordinates of the previous layout as an argument so that they can try to construct a suitably smooth sequence of node positions. They also include the `default.dist` parameter which can be tweaked to increase or decrease the spacing between isolates and disconnected components. The default value for `default.dist` is `sqrt(network.size(net))`, see `?layout.dist` for more information.

It is important to remember that there are many types of networks for which these methods will probably not produce useful visualizations. We've had the most success with networks that are fairly sparse, where a relatively small number of ties are changing between time slices, and node turnover is not too high.

5.1   Kamada-Kawai adaptation

The function `network.layout.animate.kamadakawai` is essentially a wrapper for `network.layout.kamadakawai`. It computes a symmetric geodesic distance matrix from the input network (replacing infinite values with `default.dist`), and seeds the initial coordinates for each slice with the results of the previous slice in an attempt to find solutions that are as close as possible to the previous positions. It is not as fast as MDSJ, and the layouts it produces are not as smooth. But it has the advantage of being written entirely in R, so it doesn't have the pesky external dependencies of MDSJ. For this reason it is the default layout algorithm.

5.2   MDSJ (Multidimensional Scaling for Java)

According to its authors:

> MDSJ (MDSJ , 2009) is a free Java library for Multidimensional
> Scaling (MDS). It is a free, non-graphical, self-contained, lightweight
> implementation of basic MDS algorithms and intended to be used
> both as a standalone application and as a building block in Java
> based data analysis and visualization software.

MDSJ is a very efficient implementation of MDS so `network.layout.animate.MDSJ`
gives the best performance of any of the algorithms tested so far – despite the
overhead of writing matrices out to a Java program and reading coordinates
back in. Like all of the MDS-variants, MDSJ will check and give errors if you
try to call it with a non-symmetric distance matrix. Currently `max_iter` is the
only user argument that is passed through to the Java wrapper. It controls the
maximum number of optimization steps. The default value is 50 which is usu-
ally sufficient. But it can be increased for layouts that appear to be not entirely
converging, or perhaps decreased to save some speed on simpler layouts.

Please note that the MDSJ library is released under Creative Commons
License "by-nc-sa" 3.0. This means using the algorithm for commercial purposes
would be a violation of the license. Due to CRAN's license restrictions, the
MDSJ binary is not distributed along with the (GPL-licensed) `ndtv` package.
Instead, the first time the layout is called, it will ask if you want to automatically
download and install the library. More information about the MDSJ library
and its licensing can be found at `http://www.inf.uni-konstanz.de/algo/software/mdsj/`.

5.3   Use a TEA attribute

The `useAttribute` layout is useful if you already know exactly where each
vertex should be drawn at each timestep, and you just want to render out the
network. It just needs to know the names of the dynamic attribute holding the
x coordinate and the y coordinate for each time step.

5.4   Graphviz

The Graphviz layout is a wrapper for the Graphviz `http://http://www.graphviz.org` software library (John Ellson et al , 2001). If the library is installed on
your system (see verb@?install.graphviz@), it provides a number of additional
high-quality layouts. When layout is called it checks for a working Graphviz
installation (falling back to Kamada-Kawai if Graphviz cannot be found) and
writes the network to a temp file using `export.dot`. Then the appropriate
Graphviz layout engine (default is `neato`) is executed via a `system` call, and the
coordinates of the vertices are parsed from the output.

Currently, the arguments to `layout.par` can be used to specify the Graphviz
layout engine to use (i.e. `gv.engine='neato'` for stress-minimized, `gv.engine='dot'`
for hierarchical, `gv.engine='fdp'` for force-directed, etc) and additional command-
line control parameters can be passed in via `gv.args`. For example, to use the
'dot' layout, but change layout rank direction to Left-Right:

```
> layout.par=list(gv.engine='dot',gv.args='-Grankdir=LR')
```

See `http://www.graphviz.org/content/command-line-invocation`. Note that Graphviz's graphic rendering parameters are not used to control network plot rendering (but they may impact layout positions).

5.5 User-generated layout functions

We can define new layout functions by following the appropriate naming structure. For example, if we wanted a layout that just arranged all the active vertices in a circle we could define a new function `network.layout.animate.circle`.

```
> network.layout.animate.circle <- function(net, dist.mat = NULL,
+         default.dist = NULL, seed.coords = NULL, layout.par = list(),
+         verbose=FALSE){
+
+     n<-network.size(net)
+     x<-10*cos( seq(0,2*pi, length.out=n))
+     y<-10*sin( seq(0,2*pi, length.out=n))
+     return(cbind(x,y))
+   }
```

We can then re-compute a new animation for the simulation output using our new "circle' layout function.

```
> stergm.sim.1<-compute.animation(stergm.sim.1,
+             slice.par=slice.par,animation.mode='circle')
> render.animation(stergm.sim.1)
> ani.replay()
```

5.6 Other techniques

We have tested some layouts using R libraries for doing SMACOF (de Leeuw , 2009) and standard MDS optimization. The former gave high-quality results but was extremely slow, the later often didn't give stable results. Both may be included in future releases of `ndtv` if the performance issues improve.

6 Vertex dynamics

Edges are not the only things that can change in networks. In some dynamic network data-sets vertices also enter or leave the network (become active or inactive). Lin Freeman's windsurfer social interaction data-set (Almquist et all, 2011) is a good example of this. In this data-set there are different people present on the beach on different days, and there is even a day of missing data. These networks also have a lot of isolates, which tends to scrunch up the rest of the components so they are hard to see. Setting a lower `default.dist` can help with this.

```
> data(windsurfers)
> slice.par<-list(start=1,end=31,interval=1,
+                 aggregate.dur=1,rule="latest")
> windsurfers<-compute.animation(windsurfers,slice.par=slice.par,
+                                default.dist=3,
+                                animation.mode='MDSJ',
+                                verbose=FALSE)
> render.animation(windsurfers,vertex.col="group1",
+                  edge.col="darkgray",
+                  displaylabels=TRUE,label.cex=.6,
+                  label.col="blue", verbose=FALSE)
> ani.replay()
```

In this example[6] the turnover of people on the beach is so great that structure appears to change chaotically, and it is quite hard to see what is going on. Notice also the blank period at day 25 where the network data is missing. There is also a lot of periodicity, since a lot more people go to the beach on weekends. So in this case, lets try a week-long slice by setting `aggregate.dur=7` to try to smooth it out so we can see some structure.

```
> slice.par<-list(start=0,end=24,interval=1,
+                 aggregate.dur=7,rule="latest")
> windsurfers<-compute.animation(windsurfers,slice.par=slice.par,
+                                default.dist=3,
+                                animation.mode='MDSJ',
+                                verbose=FALSE)
> render.animation(windsurfers,vertex.col="group1",
+                  edge.col="darkgray",
+                  displaylabels=TRUE,label.cex=.6,
+                  label.col="blue", verbose=FALSE)
> ani.replay()
```

This new rolling–"who interacted this week" network[7] is larger and more dense (which is to be expected) and also far more stable. There is still some turnover due to people who don't make it to the beach every week but is possible to see some of the sub-groups and the the various bridging individuals.

7 Animating graphic attributes

7.1 Using dynamic attributes (TEAs)

If a network has dynamic attributes defined, they can be used to define graphic properties of the network which change over time. We can activate some attributes on our earlier "wheel" example, setting a dynamic attribute for edge widths:

---

[6]http://statnet.csde.washington.edu/movies/ndtv_vignette/windsurfers_v1.mp4
[7]http://statnet.csde.washington.edu/movies/ndtv_vignette/windsurfers_v2.mp4

```
> activate.edge.attribute(wheel,'width',1,onset=0,terminus=3)
> activate.edge.attribute(wheel,'width',5,onset=3,terminus=7)
> activate.edge.attribute(wheel,'width',10,onset=3,terminus=Inf)
```

We must make sure the attributes are always defined for each time period that the network will be plotted or else an error will occur. So we first set a default value from `-Inf` to `Inf` before defining which elements we wanted to take a special value.

```
> activate.vertex.attribute(wheel,'mySize',1, onset=-Inf,terminus=Inf)
> activate.vertex.attribute(wheel,'mySize',3, onset=5,terminus=10,v=4:8)
```

We can set values for vertex colors.

```
> activate.vertex.attribute(wheel,'color','gray',onset=-Inf,terminus=Inf)
> activate.vertex.attribute(wheel,'color','red',onset=5,terminus=6,v=4)
> activate.vertex.attribute(wheel,'color','green',onset=6,terminus=7,v=5)
> activate.vertex.attribute(wheel,'color','blue',onset=7,terminus=8,v=6)
> activate.vertex.attribute(wheel,'color','pink',onset=8,terminus=9,v=7)
```

Finally we render it, giving the names of the dynamic attributes to be used to control the plotting parameters for edge with, vertex size, and vertex color.

```
> render.animation(wheel,edge.lwd='width',vertex.cex='mySize',
+                   vertex.col='color',verbose=FALSE)
> ani.replay()
```

The attribute values for the time points are defined using `network.collapse`, which controls the behavior if multiple values are active for the plot period.

7.2   Functional plot arguments

Sometimes it is awkward or inefficient to pre-generate dynamic attribute values. Why create and another attribute for color if it is just a simple transformation of an existing attribute or measure? The `render.animation` function has the ability to accept the `plot.network` arguments as functions with special arguments to be evaluated on the fly at each time point as the network is rendered. So, for example, if we wanted to use our previously created "width" attribute to control the color of edges along with their width:

```
> render.animation(wheel,edge.lwd=3,
+     edge.col=function(slice){rgb((slice%e%'width')/10,0,0)},
+     verbose=FALSE)
> ani.replay()
```

Notice the use of the `slice` argument to the function instead of the original name of the network. The arguments of plot control functions must draw from a specific set of named arguments which will be substituted in and evaluated at each time point before plotting. The set of valid argument names is:

17

- net is the original (un-collapsed) network

- slice is the network collapsed to be rendered with the appropriate onset and terminus

- s is the slice number in the sequence to be rendered

- onset is the onset (start time) of the slice to be rendered

- terminus is the terminus (end time) of the slice to be rendered

So in the example above, at each time point the edge attribute "width" is extracted and used to control the red component of the RGB color. We can also define functions based on network measures such as betweenness:

```
> require(sna)
> wheel%n%'slice.par'<-list(start=1,end=10,interval=1,
+                           aggregate.dur=1,rule='latest')
> render.animation(wheel,
+       vertex.cex=function(slice){(betweenness(slice)+1)/5},
+       verbose=FALSE)
> ani.replay()
```

In this example we had to modify the start time using the slice.par setting to avoid time 0 because the betweenness function will give an error for a network with no edges. The main plot commands accept functions as well, so it is possible to do fun things like implement a crude zoom effect by setting xlim and ylim parameters to be dependent on the time.

```
> render.animation(wheel,
+       xlim=function(onset){c(-5/(onset*.5),5/(onset*.5))},
+       ylim=function(onset){c(-5/(onset*.5),5/(onset*.5))},
+       verbose=FALSE)
> ani.replay()
```

8   Exploring proximity with timelines

We've already introduced the timeline function in the section of slicing time. Although it can be helpful for debugging and revealing the over density of events in a dynamic network, it is difficult to understand what those events imply for changes in network structure and connectivity. The movies do a reasonably good job (at least for sparse networks) of illustrating the moment-to-moment changes in structure, but it is often hard to grasp the overall shifts without rewinding and replaying the movie over-and-over.

The proximity.timeline attempts a sort of compromise. It collapses all the momentary structure information down to a single vertical dimension, and uses the horizontal axis for time. More precisely, the network is extracted at

each time bin and the geodesic distances are computed. But instead of creating a 2-dimensional layout as `compute.animation` does, the network layout is a single dimension indicating how relatively 'close' or 'far' the vertices are from each other. Each vertex's positions in the the time steps are linked together by a spline. So, like the timeline, each vertex traces out a horizontal trajectory, but in this case it can swerve diagonally up and down as it moves from group to group.

For example, if we return to the Stergm simulation example, we can contrast the (entirely fictional, simulated) histories of the Tornabouni and Strozzi family marriage alliances as a blue and green lines on the proximity timeline. First we load in a short example dataset of the flomarriage simulation.

```
> data(short.stergm.sim)  # load a short example dataset of the flomarriage simulation
> proximity.timeline(short.stergm.sim,mode='sammon',
+                    default.dist=10,
+                    labels.at=c(1,16,25),
+                    label.cex=0.7,
+                    vertex.col=c(rep('gray',14),'green','blue') )
```

Initially, Tornabouni is part of the component many of the other families. Around t=9, they split off and become isolated, but then pair up with Laberteschi at t=21, and then rejoin one of the groups from the big component (which itself split in half at t=19). For most of the simulation the Strozzi (green), trace out a relatively horizontal existence as an isolate, but eventually connect with the Acciaiuoli near the end.

Perhaps another way to illustrate how the `proximity.timeline` works is to combine it with several static snapshots of the network (like those produced by `filmstrip`).

The `toy_epi_sim` dataset is an example network of a trivial simulated disease process spreading over a simulated dynamic contact network among 100 individuals for 25 discrete time steps. It was produced by the EpiModel package, and it includes an attribute named 'ndtvcol' corresponding to the simulated infection status of the vertices. The infection status changes over time.

```
> data(toy_epi_sim)
> # set up layout to draw plots under timeline
```

```
> layout(matrix(c(1,1,1,2,3,4),nrow=2,ncol=3,byrow=TRUE))
> # plot a proximity.timeline illustrating infection spread
> proximity.timeline(toy_epi_sim,vertex.col = 'ndtvcol',
+                    spline.style='color.attribute',
+                    mode = 'sammon',default.dist=100,
+                    chain.direction='reverse')
> # plot 3 static cross-sectional networks
> # (beginning, middle and end) underneath for comparison
> plot(network.collapse(toy_epi_sim,at=1),vertex.col='ndtvcol',
+      vertex.cex=2,main='toy_epi_sim network at t=1')
> plot(network.collapse(toy_epi_sim,at=17),vertex.col='ndtvcol',
+      vertex.cex=2,main='toy_epi_sim network at=17')
> plot(network.collapse(toy_epi_sim,at=25),vertex.col='ndtvcol',
+      vertex.cex=2,main='toy_epi_sim network at t=25')
> layout(1)
```


This plot can be hard to read as a small image, it is worth rendering it in a

big plot window. In the first network snapshot we see a few scattered infections (in red), some small components and a medium size component. These groups show up in the the beginning of the timeline as bundles of lines, with larger bundles corresponding to the larger components. There are several red lines indicating the infections mixed in with the blue threads. As time progresses, the bundles untwist and braid as vertices split off to join other components. Some vertices become isolates and tend to fly off to the top and bottom of the chart.

The number of red lines grows as the infection spreads. By t=17 (the second snapshot) the network has briefly formed a large component, viable in the timeline as a fat bundle in the center. At the end of the simulation (t=25) most of the network has become infected, and the large component has broken up again into multiple medium-sized components.

Of course there are quite a few vertices (100) so it it is difficult to see exactly what is going on in detail, especially when they cross over each other. But the proximity timeline is sometimes able to illustrate the features of the forward reachable paths and changes in overall network structure in ways that can be missed when viewing a movie.

Naturally the proximity.timeline plots suffer from some of the same noise and reproduceability problems that challenge the network layouts. The geodesic distance information is tightly compressed onto a single dimension, so the exact ordering of the vertices in any specific region may be due to chance, just as the rotation and relative positions of components in a network plot are not directly meaningful. This is also a fairly experimental tool, so getting good results still requires playing around a bit with the various algorithms and adjusting `default.dist` to a value larger enough to force clusters close enough together without making them overlap too much.

8.1   Examining missingness

This plot shows the timeline view vertex activity spells for the "windsurfer" dataset.

```
> timeline(windsurfers,plot.edge.spells = FALSE)
```

Notice that in contrast to the McFarland classroom dataset we saw earlier, many of the vertices are almost entirely unobserved.

```
> proximity.timeline(windsurfers,mode='sammon')
```

23

The proximity.timeline function uses the spline.style argument to control how the spline segments corresponding to vertex inactivity should be rendered.

We can also the use of the start and end paramter to zoom in and render only a portion of the time range.

```
> proximity.timeline(windsurfers,start=20,end=31,mode='sammon',
+                    spline.style='inactive.gaps')
```

```
> proximity.timeline(windsurfers,start=20,end=31,
+                    mode='sammon',spline.style='inactive.ignore')
```

approx. distance among vertices of windsurfers — time

```
> proximity.timeline(windsurfers,start=20,end=31,mode='sammon',
+                    spline.style='inactive.ghost')
>
```

26

The ghost version (which is the default) offers a compromise by linking the observed spells with a faint dotted line, making it possible to trace the trajectories across time without it appearing that there is much more data availible than is actually the case.

9 Dependencies for Animations

9.1 Java (for MDSJ)

In order to use the MDSJ layout algorithm, you must have Java installed on your system. Java should be already installed by default on most Mac and Linux systems. If it is not installed, you can download it from `http://www.java.com/en/download/index.jsp`. On Windows, you may need to edit your 'Path' environment variable to make Java executable from the command-line.

9.2 FFmpeg

FFmpeg `http://ffmpg.org` is a cross-platform tool for converting and rendering video content in various formats. It is used as an external library by the `animation` package to save out the animation as a movie file on disk. (see `?saveVideo` for more information.) Since FFmpeg is not part of R, you will need to install it separately on your system for the save video functionality to work. The instructions for how to do this will be different on each platform. You can also access these instructions using `?install.ffmpeg`

10 Compressing video

The saved video output of the animation often produces very large files. These may cause problems for your viewers if you upload them directly to the web. It is almost always a good idea to compress the video, as a dramatically smaller file can usually be created with little or no loss of quality. Although it may be possible to give `saveVideo()` various `other.opts` to control video compression[8], determining the right settings can be a trial and error process. Handbrake `http://handbrake.fr/` is an excellent and easy to use tool for doing video compression into the web-standard H.264 codec with appropriate presets.

11 Reference for the main commands

Included here are more complete explanations of the main function. You can also refer to the man pages `?compute.animation` and `?render.animation`.

11.1 compute.animation()

The `compute.animation()` function computes a sequence of vertex layouts suitable for rendering a network animation. It steps through a `networkDynamic` object and applies layout algorithms at specified intervals, storing the calculated coordinates in the network for later use by the `render.animation` function. Generally the layouts are done in a sequence with each using the previously calculated positions as initial seed coordinates in order to smooth out the resulting movie.

The command takes several important arguments as named elements of the `slice.par` list. The parameters indicate it how "slice up' the network when computing layouts (`start`, `end`, `aggregate.dur` and `rule`), what type of layout algorithm to use (`animation.mode`), possible parameters to control the layouts (as a list named `layout.par`) and how much to try to separate nodes or disconnected components (`default.dist`). The computed coordinates are stored as dynamic vertex attributes named `animation.x.active` and `animation.y.active`. The

---

[8]The default settings for ffmpeg differ quite a bit depending on platform, some installations may give decent compression without tweaking the settings

slice `slice.par` list is stored as a standard network attribute. The network argument is modified in place, and returned invisibly

For each time slice, new coordinates are only computed for the active set of vertices, so the function usually behaves appropriately for networks with changing vertex sets.

The other parameters are as follows

- `seed.coords` an (optional) array of initial coordinates to be used for the very first layout in the sequence or when vertices first pop into existance.

- `weight.attr` can provide the name of a numeric edge attribute defining weights for edges to be interpreted by the layout algorithm. The values `activity.duration` or `activity.count` can be used to weight edges by the duration or count of the edge's activity spells in the time slice. the `weight.dist` parameter determines if the weights should be treated as similarities (larger values means closer vertices) or distances.

- `chain.direction` a value of 'forward' indicates the chain of layouts should be computes in forward temporal order. A value 'reverse' runs the chain backwards. For some layouts, reverse-chaining means that isolated vertices are more likely to have positions close to the partners they will be tied to.

11.2   render.animation()

This function is designed to step through a network object extracting slice networks according to the previously cached `slice.par` settings. It retrieves the `animation.x` and `animation.y` coordinates for each slice and passes them to `plot.network` to render the frame. If no `slice.par` network attribute is found to define the time range to render it will make one up using the smallest and largest non-Inf time values and unit-length non-overlapping time steps. If no stored coordinates are found it will call `compute.animation`. Additional `plot.network` control parameters (to set colors, line widths, etc) can be passed in via the `...` arguments. See `?plot.network` for the full list.

As mentioned earlier, a number of "tweening" animation frames are generated between each network slice with the positions of the vertices interpolated between the slices. This creates the illusion of smooth motion as the vertices change position, making it much easier to visually track changes in the network structure. As each slice (and tweening slice) is plotted, `ani.record` is called to store the image as a frame of the animation for later output.

Parameters to control the animation are read from a list passed in via the `render.par` argument.

- `tween.frames` is the number of interpolated frames to generate between each pre-calculated network layout. Default is 10. Increasing this will make the animation appear smoother and slower, but will make the file sizes much larger.

- `show.time` defaults to `TRUE`, in which case the x-axis of the plot will be labeled with the onset and terminus time for each slice as it is shown.

- `show.stats` does nothing with its default value of `NULL`. But if it is set to a string, it is assumed to be a formula and will passed to `summary.stergm` and the results used to display the network statistics for the current slice on the plot.

- `extraPlotCmds` provides a way to present additional information (such as annotations) on the plot. The value of this argument will be passed to `eval()` after each frame has been plotted, so drawing commands can be added here.

There are also several lists of arguments that give default values that will be passed to the appropriate lower-level commands. The `plot.par` list is passed to the `par()` command and provides a way to configure some of the general plot details such as background color, margins, fonts, etc. Similarly, the `ani.options` list is passed to the `ani.options()` command to configure settings for the animation package such as `interval` to control the time between frames in playback.

The `render.cache` argument provides a way to control the caching of the plot frames. The default value of verb@render.cache='plot.list'@ causes each frame of the animation to be stored in an internal list by the `ani.record` function of the animation library. This is very useful for testing and replaying animations in R's plot window, but can be very slow (or cause out-of-memory errors) for large animations. If the value is set to verb@render.cache='none'@, the plot will not be recorded (but can be saved directly to disk via `saveVideo()`) and cannot be replayed via the `ani.replay()` function.

11.3 saveVideo()

The `animation` package provides several neat tools for storing animations once they have been rendered.

- `ani.replay()` plays the animation back in the R plot window. (see `?ani.options` for more parameters)

- `saveVideo()` saves the animation as a movie file on disk (if the FFmpeg library is installed).

- `saveGIF()` creates an animated GIF (if ImageMagick installed)

- `saveLatex()` creates an animation embedded in a pdf (didn't work for me...)

Please see `?animation` and each function's help files for more details. With the exception of `ani.replay()` each of these requires the presence of some external library software which may need to be installed on your system as described in Dependencies (section 9).

11.4 render.d3movie

The `render.d3movie` can save out a network animation an interactive HTML5 SVG to display in a web browser. Animations are generated using a process nearly identical to `render.animation`. However, instead of using R's plotting functions and the animation library, the relevant information is cached and written into a JSON-formatted file, embedded into a web page along with ndtv-d3 player, and displayed in a web browser. Details and additional examples are included in the ndtv-d3 vignette: `vignette('ndtv-d3 animation examples',package='ndtv')`

12 Limitations

12.1 Size limits

Like most network algorithms, the layouts tend to scale quite badly. We generally have only had enough patience to generate moves for networks of less than 1000 vertices. There also seems to be quite a bit of overhead in the `animation` package, so the generation process seems to slow down considerably for longer duration networks or when slice or render parameters cause lots of slices to be generated.

References

Algorithmics Group, University of Konstanz (2009) *MDSJ: Java Library for Multidimensional Scaling (Version 0.2).* `http://www.inf.uni-konstanz.de/algo/software/mdsj/`.

Almquist, Zack W. and Butts, Carter T. (2011). "Logistic Network Regression for Scalable Analysis of Networks with Joint Edge/Vertex Dynamics." *IMBS Technical Report MBS 11-03*, University of California, Irvine.

Bender-deMoll, Skye and McFarland, Daniel A. (2006) The Art and Science of Dynamic Network Visualization. *Journal of Social Structure. Volume 7, Number 2* `http://www.cmu.edu/joss/content/articles/volume7/deMollMcFarland/`

Bender-deMoll, S., Morris, M. and Moody, J. (2008) Prototype Packages for Managing and Animating Longitudinal Network Data: dynamicnetwork and rSoNIA *Journal of Statistical Software* 24:7.

Butts CT (2008). network: A Package for Managing Relational Data in R. *Journal of Statistical Software*, 24(2). `http://www.jstatsoft.org/v24/i02/`.

Butts C, Leslie-Cook A, Krivitsky P and Bender-deMoll S (2015). *networkDynamic: Dynamic Extensions for Network Objects*. R package version 0.8, `http://statnet.org`.

de Leeuw J and Mair P (2009). "Multidimensional Scaling Using Majorization: SMACOF in R." *Journal of Statistical Software*, **31**(3), pp. 1–30. `http://www.jstatsoft.org/v31/i03/`

Bender-deMoll S (2015). *ndtv: Network Dynamic Temporal Visualizations*. R package version 0.6, `http://statnet.org`.

John Ellson et al (2001) Graphviz – open source graph drawing tools *Lecture Notes in Computer Science*. Springer-Verlag. p483-484 `http://www.graphviz.org`

Handcock MS, Hunter DR, Butts CT, Goodreau SM, Morris M (2003b). statnet: Software tools for the Statistical Modeling of Network Data. Statnet Project, Seattle, WA. Version 3, `http://www.statnetproject.org.`

Hunter DR, Handcock MS, Butts CT, Goodreau SM, Morris M (2008b). ergm: A Package to Fit, Simulate and Diagnose Exponential-Family Models for Networks. *Journal of Statistical Software*, 24(3). `http://www.jstatsoft.org/v24/i03/.`

Krivitsky P and Handcock M (2014). *tergm: Fit, Simulate and Diagnose Models for Network Evolution based on Exponential-Family Random Graph Models*. The Statnet Project (`http://www.statnet.org`). R package version 3.2.4, `CRAN.R-project.org/package=tergm`.

McFarland, Daniel A. (2001) "Student Resistance: How the Formal and Informal Organization of Classrooms Facilitate Everyday Forms of Student Defiance." *American Journal of Sociology* **107** (3): 612-78.

Greg Michalec, Skye Bender-deMoll, Martina Morris (2014) "ndtv-d3: an HTML5 network animation player for the ndtv package" The statnet project. http://statnet.org

Newcomb T. (1961) *The acquaintance process* New York: Holt, Reinhard and Winston.

Xie Y (2013). "animation: An R Package for Creating Animations and Demonstrating Statistical Methods." *Journal of Statistical Software*, **53**(1), pp. 1–27. `http://www.jstatsoft.org/v53/i01/.`